

The background features a large, light blue watermark logo of the Technical University of Cluj-Napoca. The logo consists of a shield with a stylized 'T' and 'U' inside, with the text 'TECHNICAL UNIVERSITY' at the top, 'OF CLUJ-NAPOCA' in the middle, and 'Computer Science' at the bottom.

Fundamental Algorithms

Lecture #7

Cluj-Napoca

Computer Science

Agenda

- **Augmented trees**
 - **Type 2 – Tree/lists**
- **Balanced trees – why and how (review)**
- **Red-Black Trees (balanced trees type 3)**

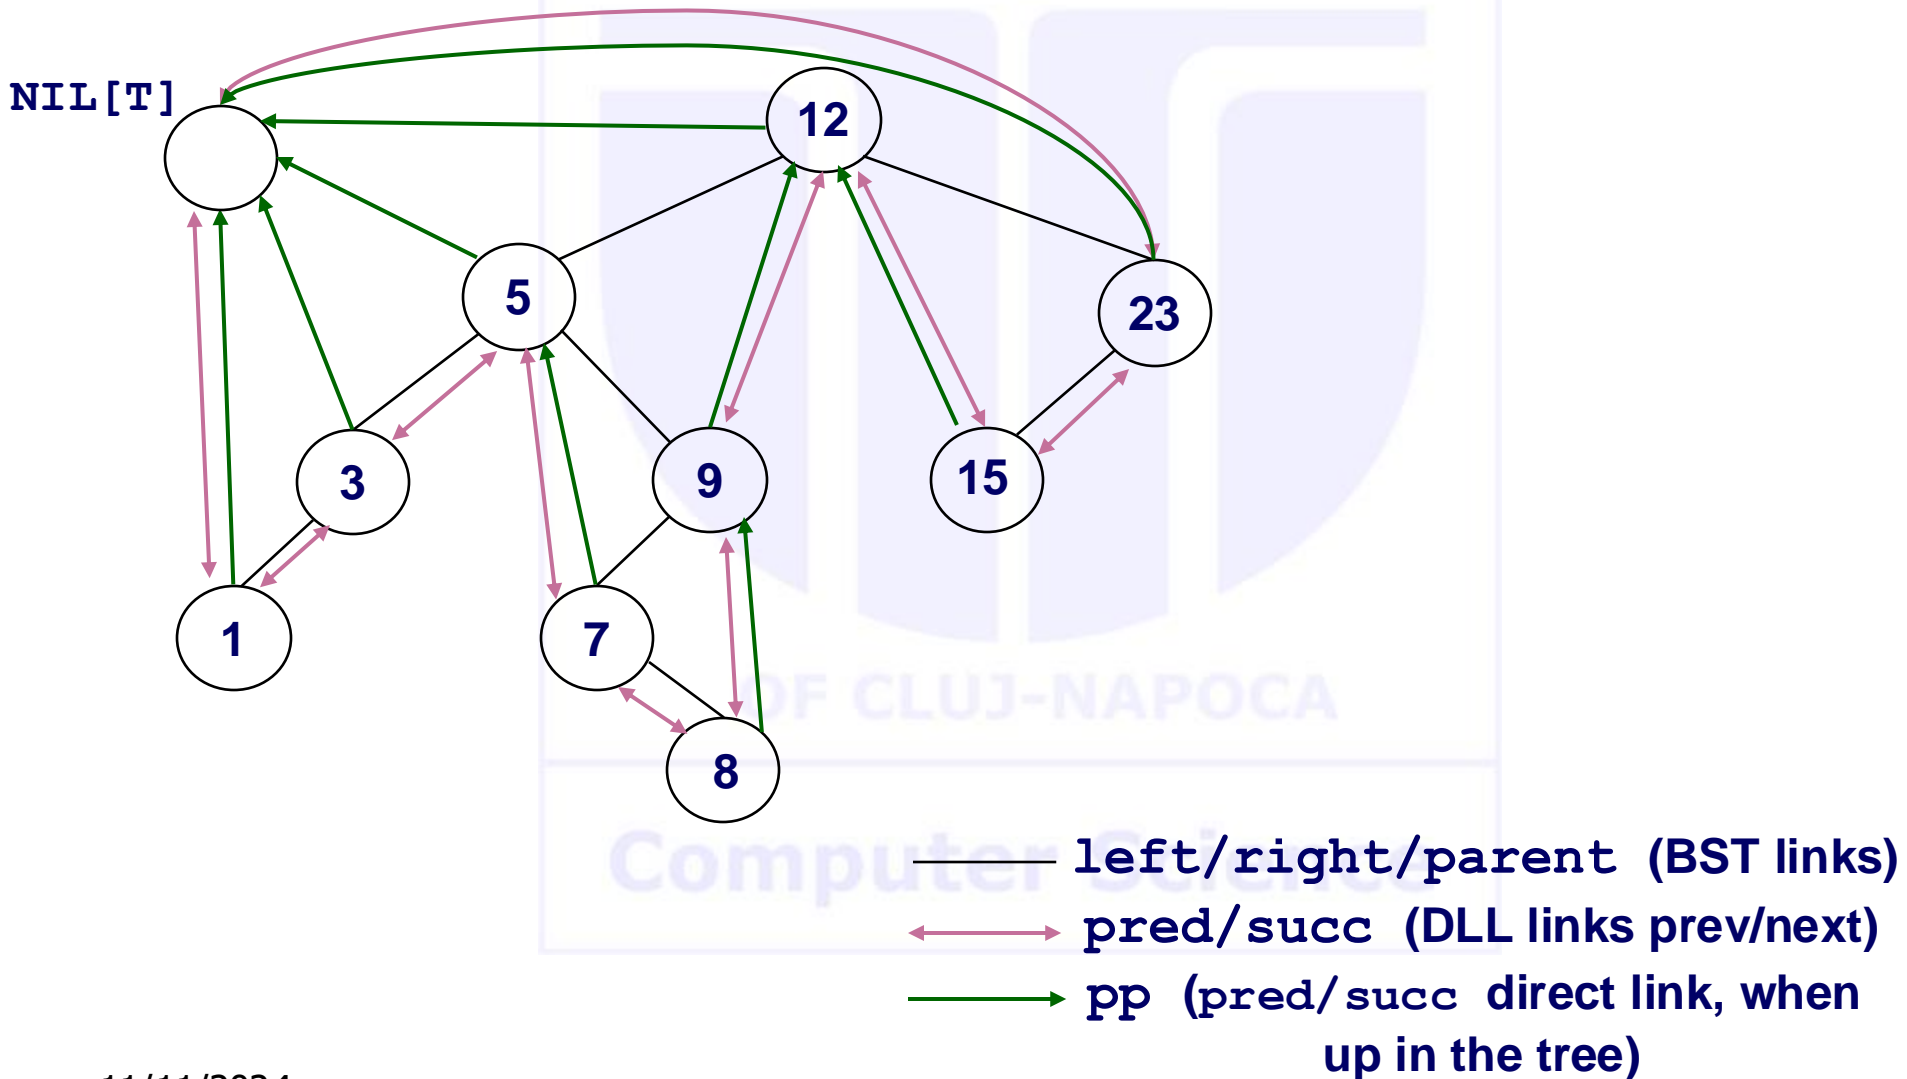
Augmented trees (type 2)

- Requirements:
 - Regular operations are performed as (**same performance also**) in BST (walk ($O(n)$), search, ins, del ($O(h)$))
 - Several other operations are enhanced (i.e. performe faster)
 - Succ
 - Pred
 - Min
 - Max
 - **All required to be performed in $O(1)$!!!**
 - BUT NONE of the before-defined operations should degrade their performance

Augmented trees – contd.

- Info in a node:
 - Usual info:
 - key
 - left pointer
 - right pointer
 - parent pointer
 - Supplementary info (see picture on the blackboard):
 - succ pointer
 - pred pointer (together ensure walking through the list)
 - pp ensures min/max oper. (in a regular BST, succ/pred calculated **either** based on min/max **or** pp - which is determined at the execution time)

Example



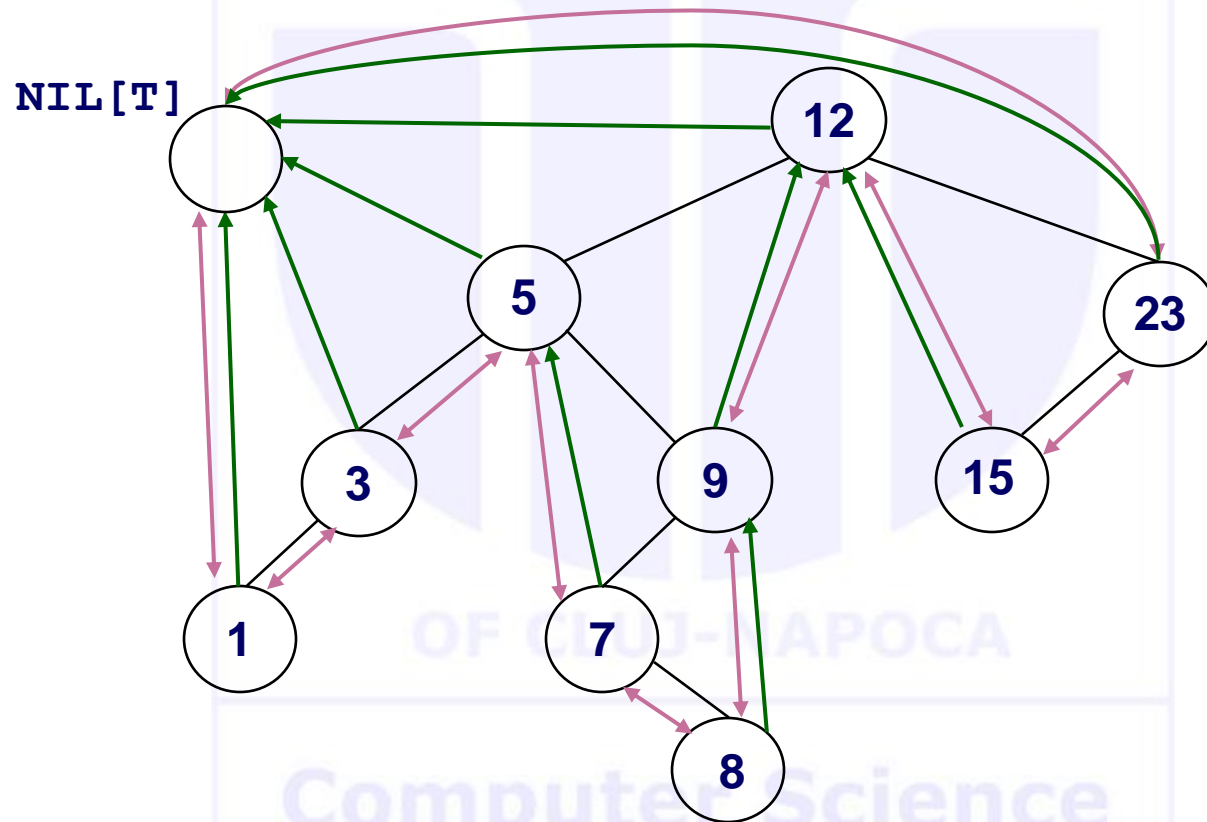
Augmented trees – contd.

- The structure acts BOTH as a BST and DLL!!
- Regular operations are:
 - done like in any other BST
 - in addition, need to make some updates
- They (the additional updates) refer to:
 - making the appropriate links within the DLL (set/update the *pred* and *succ* pointers)
 - link the double pointer (set/update the *pp* pointer)

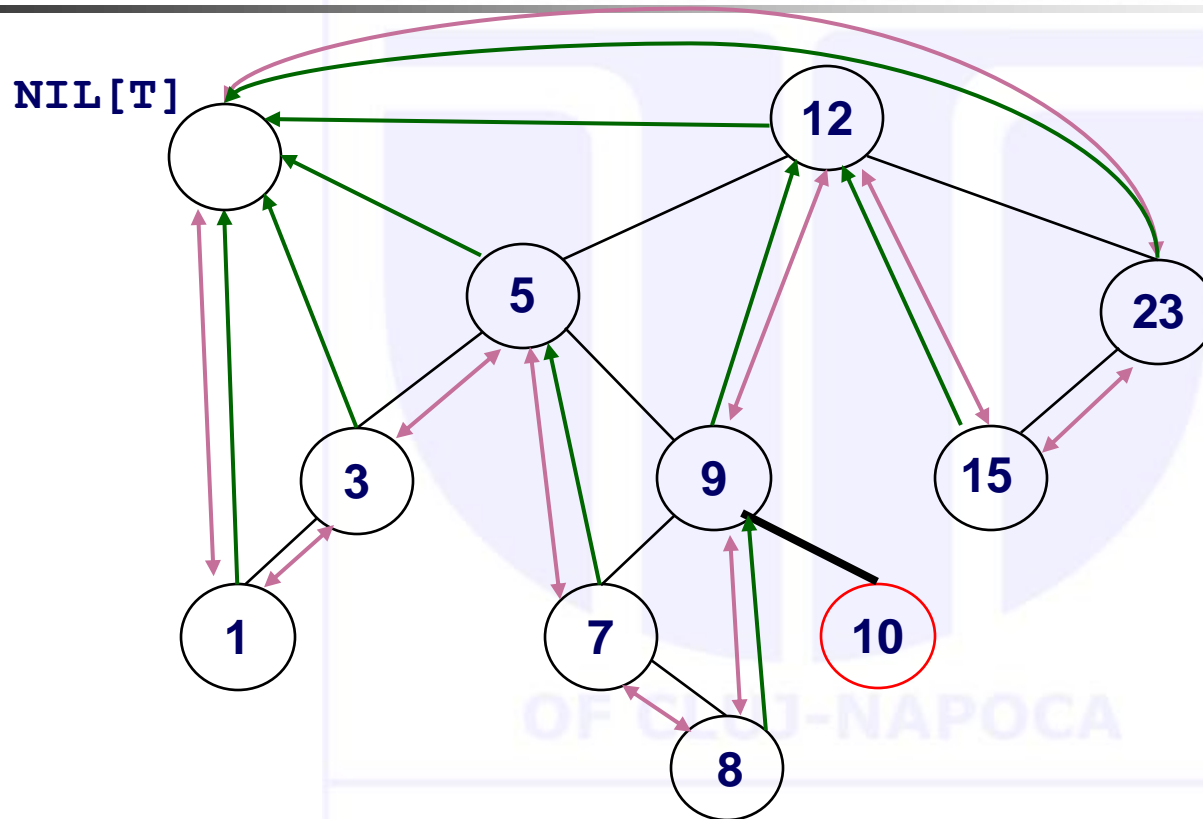
Augmented trees – Insert

- Regular **insert** operation in a BST (x inserted) **+**
if $x = \text{right}[p[x]]$ //node inserted = right child
then //case #1
 $pp[x] \leftarrow \text{succ}[p[x]]$
 $\text{dl_list_ins_after}(p[x], x)$
else //case #2; node inserted = left child
 $pp[x] \leftarrow \text{pred}[p[x]]$
 $\text{dl_list_ins_after}(pp[x], x)$

AugBST-Insert(T, 10)




AugBST-Insert(T, 10)




if $x = \text{right}[p[x]]$ //node inserted = right child
then //case #1

$pp[x] \leftarrow \text{succ}[p[x]]$
 $dl_list_ins_after(p[x], x)$



11/11/2024

AugBST-Insert(T, 10)



```

if  $x = \text{right}[p[x]]$  //node inserted = right child
    then //case #1
         $pp[x] \leftarrow \text{succ}[p[x]]$ 
         $dl\_list\_ins\_after(p[x], x)$ 
    
```

Augmented trees – Delete

(z = node requested to be removed; it's content is replaced by y's content
y=node actually removed = at most 1 child node;
x = its (y) only child/if any, might be nil;
z=y if z has at most one child)

- Apply regular **delete** operation in a BST + code below

if right[y]=nil

then

x<-left[y]

while x<>nil

do pp[x]<-pp[y] //update pp


x<-right[x]

dl_list_del(y)

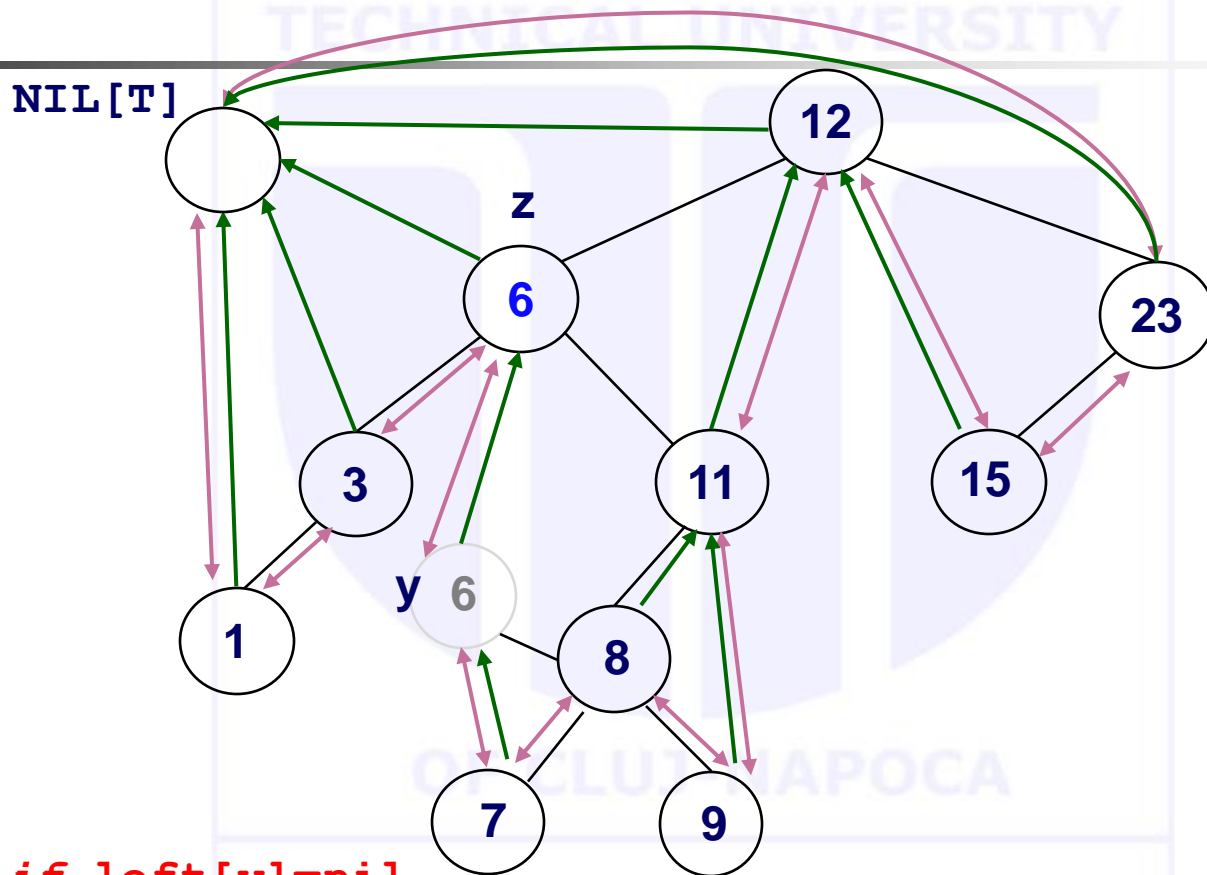
else

//symmetric on the left

BST-Delete(T,5)



BST-Delete(T,5)



else if left[y]=nil

then

x<-right[y]

while x <> nil

do

pp[x]<-pp[y]

x<-left[x]

//no child to the left

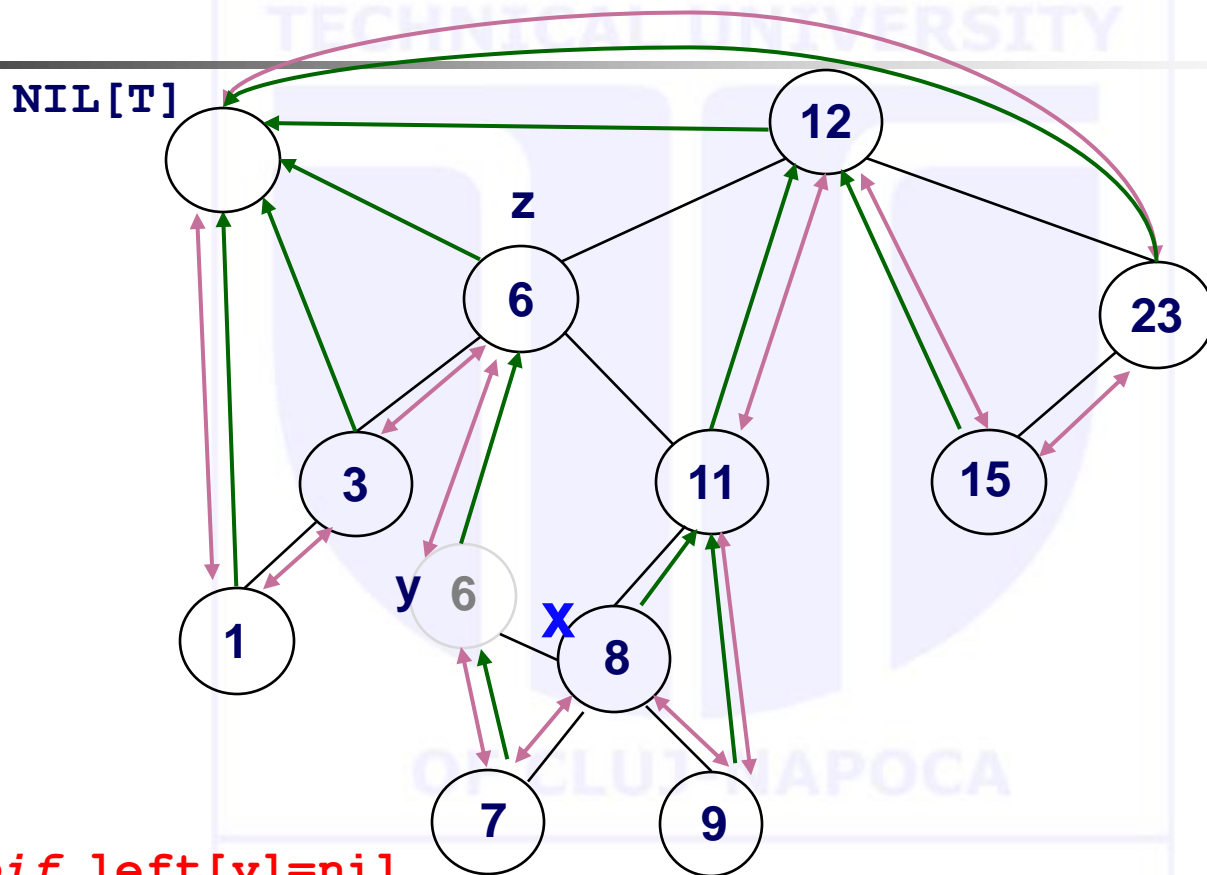
//x = y's only child

//along x's left branch

//update pp

dl_list_del(y)

BST-Delete(T,5)



elseif left[y]=nil

then

x<-right[y]

while x <> nil

do

pp[x]<-pp[y]

x<-left[x]

//no child to the left

//x = y's only child

//along x's left branch

//update pp

dl_list_del(y)



```
x<-right[y]
```

//no child to the left

//x = y's only child

//along x's left branch

```
pp[x] <- pp[y]
```

```
//update pp
```

```
x<-left[x]
```

```
dl list del(y)
```




```
x<-right[y]
```

//no child to the left

//x = y's only child

//along x's left branch

```
//update pp
```

```
x<-left[x]
```

```
dl_list_del(y)
```



```
x<-right[y]
```

//no child to the left

//x = y's only child

//along x's left branch

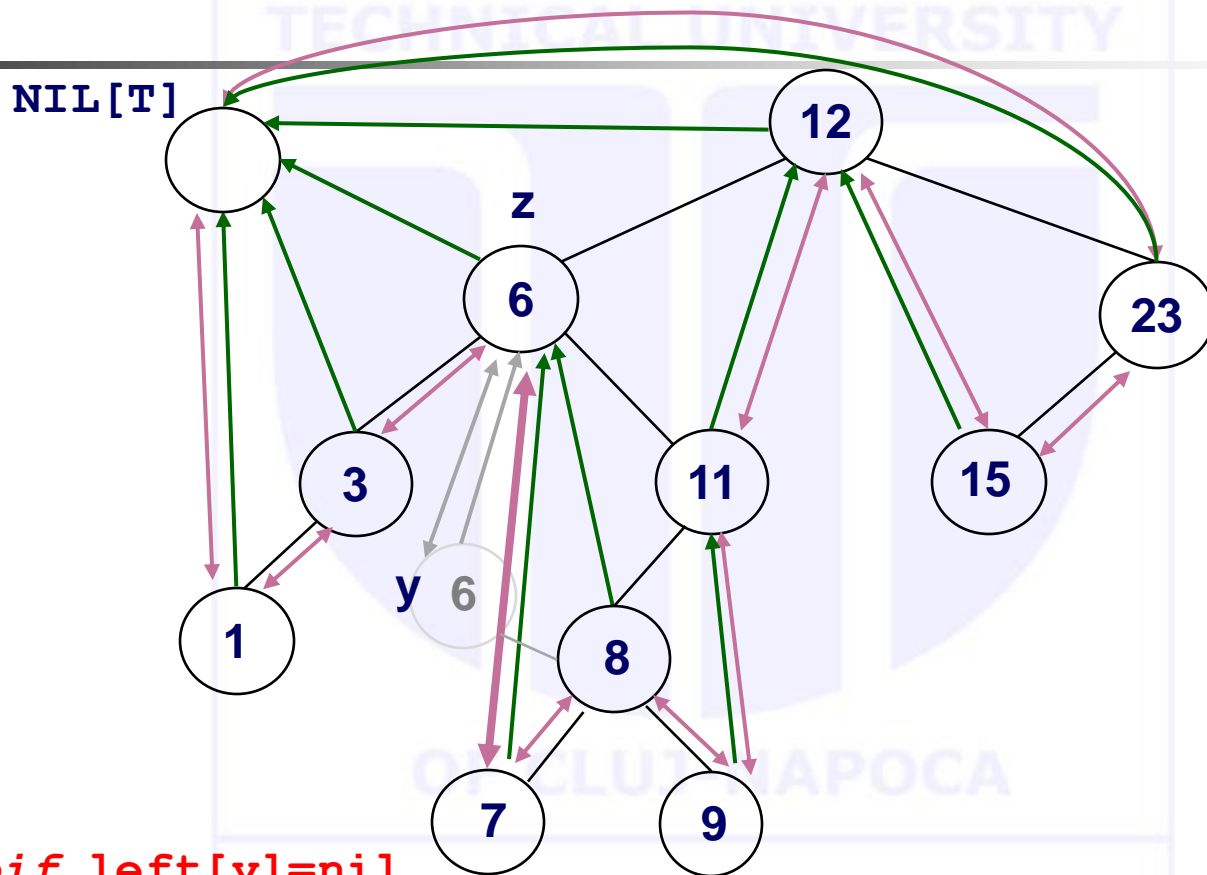
```
pp[x] <- pp[y]
```

```
//update pp
```

```
x<-left[x]
```

```
dl_list_del(y)
```

BST-Delete(T,5)



elseif left[y]=nil

then

x←right[y]

while x <> nil

do

pp[x]←pp[y]

x←left[x]

//no child to the left

//x = y's only child

//along x's left branch

//update pp

Augmented trees – Min

- **min** (based on *succ* and *pp* as opposed to regular BST, where *succ* is calculated based on *min* **or** determined *pp*)

if $x = \text{left}[p[x]]$

then


return $\text{succ}[pp[x]]$

//on the leftmost branch, **HAS TO BE** $pp[x] = \text{nil}!!!$

else

return $\text{succ}[p[x]]$


BST-Min(nodeX)



```
if  $x = \text{left}[p[x]]$   
then  
     $\text{return succ}[pp[x]]$ 
```

```
else  
     $\text{return succ}[p[x]]$ 
```

BST-Min(nodeX)



```
if x=left[p[x]]
  then
    return succ[pp[x]]]
```

```
else
  return succ[p[x]]]
```

Augmented trees – Max

- **max** (based on *pred* and *pp* as opposed to regular BST where *pred* is calculated based on *max* or determined *pp*)

if $x = \text{left}[p[x]]$

then

return $\text{pred}[p[x]]$

else

return $\text{pred}[pp[x]]$

//on the rightmost branch, **HAS TO BE** $pp[x] = \text{nil}$ and
 $\text{pred}[\text{nil}[T]] = \text{last node in inorder} = \text{last node in the list}$

Augmented trees – contd.

- Particular (initial) cases discussion on the blackboard!
- First **insert** (in the empty tree)

Tree_ins(*T*, *x*)

if *x* = root[*T*]

then //the node just inserted is the root = tree was empty before

 root[*T*] <- *x*

p[*x*] <- nil[*T*]

pp[*x*] <- nil[*T*]

dl_list_ins_after(*pp*[*x*], *x*)

else //the regular case described earlier

 ...

Homework: updates for delete!

Red-Black trees

- Balanced trees
- Both insert/delete operations take $O(\lg n)$, with at most $O(\lg n)$ for rebalancing

Def: A RBT is a BST with the following properties:

P₀: the root is black

P₁: each *node* is colored either **black** or **red**

P₂: each *leaf* (*NIL*) is **black**

P₃: both *children of a red node* are **black**

P₄: every *path* from any node to a leaf has the *same number of black nodes* ("black height")

RB Tree – Example

Is it a RBT? Why not?




Image taken from: <https://www.codesdope.com/course/data-structures-red-black-trees/>

Red-Black trees

Th: The height of a RB tree with n internal nodes is at most $2\lg(n+1)$

Proof: Let's denote by $bh(x)$ = the black height (without x) of node x

Step 1: Define the statement $P(bh)$ as follow:

$P(bh)$: $\forall x \in RBT$, the tree rooted by x has at least $2^{bh(x)} - 1$ nodes

Induction:

$P(0)$ $2^0 - 1 = 0$

Assume $P(bh)$ true $\Rightarrow P(bh+1)$ true?

x has 2 children; each child has the black height: if x is red: $bh(x)$
if x is black: $bh(x) - 1$

nb of internal nodes of x = nb of internal nodes of children(x) + 1

(itself) \Rightarrow at least $(2^{bh(x)-1} - 1) + (2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ *q.e.d.* (end of **Step 1**)

Red-Black trees

Step 2:

We know $P(bh)$ is true, i.e.

$P(bh)$: $\forall x \in \text{RBT}$, the tree rooted by x has at least $2^{bh(x)} - 1$ nodes (1)

By P_3 of RBT def (use contradiction to prove) $bh(x) \geq h/2$ (2)

//since after each red node comes a black one

$$\Rightarrow n \geq 2^{bh(x)} - 1 \quad (\text{from (1)})$$


$$\geq 2^{h/2} - 1 \quad (\text{from (2)})$$

$$n \geq 2^{h/2} - 1 \Leftrightarrow n+1 \geq 2^{h/2} \Leftrightarrow h/2 \leq \lg(n+1) \Leftrightarrow h \leq 2\lg(n+1)$$

q.e.d. (end of **Th proof**)


Red-Black trees - rotations

Similar to single rotations for AVL
They are symmetric



Picture from wiki

Red-Black trees - rotations



left_rotate(T, x)

//x root of rotation (points on P)

y<-right[x]

//y saves Q

right[x]<-left[y]

//right of P goes on B

if left[y]<>nil

//if B exists = is not nil

then p[left[y]]<-x //B's parent becomes P

p[y]<-p[x]

//Q's parent what was P's parent

if p[y]=nil

//P used to be the root of the tree

then root[T]<-y

else if x=left[p[x]] // the parent of P becomes the parent of Q

then left[p[x]]<-y

else right[p[x]]<-y

left[y]<-x

//P goes the left child of Q

p[x]<-y

//Q becomes the parent of P

RB-insert

- **Insert** like in ANY other BST
 - As a LEAF, as for any other BST
- **Assign** it a color
 - **RED**
- Check the properties
- Re-balance if needed (**RB-INSERT-FIXUP** – check the textbook for the complete code)
- **P₃**: both children of a **red** node are **black**
- True for the children (NILs) of the inserted node
- Not true for the inserted node, in case its parent is **RED** colored
- Cases to analyze and remove inconsistencies

RB-insert- Case#1

- **B inserted node (pointed by x)**
- Parent (A)=**RED**, uncle (D)=**RED**, grandparent (C)=**BLACK**
- $\alpha, \beta, \gamma, \delta, \varepsilon$ are RB trees (β, γ empty at first)



=> **Swap colors** between grandparent (C) and parent/uncle

parent<-**black** //no more P3 conflict A-B

uncle<-**black** //to preserve P4


grandparent<-**red** //to preserve P4

x<-grandparent



RB-insert- Case#1-eval

- P_3 may still be invalid, for the new x (i.e. C)
- Problem transferred 2 levels up in the tree (now β , γ not empty any longer)
- It takes (in the worst case) $O(h)$ to rebalance ($2\lg(n+1)/2 = \lg n$)



RB-insert- Case#2



- **B inserted node (pointed by x)**
- Parent(**A**)=**RED**, uncle (δ 's root)=**BLACK** (here is the difference compared to case #1), grandparent (C)=**BLACK**
- $\alpha, \beta, \gamma, \delta$ are RB trees; δ 's root is **BLACK**

=>

left_rotate(p(x))


//no more P3 conflict B-parent conflict

x <- left(x)



RB-insert- Case#2-eval

- Case #2 takes just $O(1)$ to apply, but
- P_3 is still invalid, for the new x (i.e. A-B conflict)
- \Rightarrow it is followed by case #3



RB-insert- Case#3


- Either **Inserted A**, or coming from #2 (node pointed by **x**)
- Parent (**B**)=**RED**, uncle (γ 's root)= **BLACK**, grandparent (**C**)=**BLACK**
- $\alpha, \beta, \gamma, \delta$ are RB trees

=>

parent <- *black*

grandparent <- *red*

right_rotate(*p[p[x]]*)



RB-insert- Case#3-eval

- Problem solved
- Each individual case takes $O(1)$
- Case #1 may repeat (up in the tree)
- Case #2 is followed by #3
- Case #3 solves the problem

RB-insert – Rebalancing eval

- Case #1 repeats up to the root $O(h)$
- Case #2+#3 \Rightarrow problem fixed $O(1)$
- Case #3 \Rightarrow problem fixed $O(1)$
- Insert $O(\lg n)$ + rebalancing
 - Worst case: #1 repeats $O(\lg n)$
 - Best case: #3 \Rightarrow 1 rotation $O(1)$
 - Other cases: #2+3 \Rightarrow 2 rotations $O(1)$
- $O(\lg n)$ overall worst time (case 1 repeats),
at most 2 rotations (case 2)

RB-delete

- Del as in regular BST + properties check to rebalance, if needed (RB-DELETE-FIXUP – check the textbook for the code)
- P4 (black height) is an issue

rb_delete(T, z)

tree_delete(T, z)

if color[y]=black

then rb_del_fix(T, x)

z=node to be removed (see picture on the blackboard)

y=node actually removed ($y \equiv z$ in case z has at most 1 child);
info in y is placed in z 's node

x=y's only child before the delete process takes place (could be nil, in case y has no children). After y is deleted, x becomes the child of y 's parent (thus, x 's parent could have now both children, one being x)

w=x's brother (after delete operation takes place; it's y 's brother before the deletion)



RB-delete

On x's branch check P4 property
x is y's (the removed node) only child

```
if color[x]=red  
    then color[x]<-black  
        //problem fixed; DONE!  
        //x brings its former father color  
    else color[x]<-double_black  
        //P1 property issue!
```

Computer Science

RB-delete- Case#1





- Issue at node **A** (pointed by x) which is **double black**!
 - A= was the only child of the deleted node
 - **Parent (B) =Black, brother (D) =Red**
 - $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees
- => B<->D color interchange +left rotate=> case 2 or 3 or 4

parent[x]<-red

brother[x]<-black //colors interchanged

left_rotate(p[x])

RB-delete- Case#2



Parent (B)=Red (if after case #1) or Black, brother (D)=Black, node C is Black

brother[x] ← red

if p[x]=red



then p[x] ← black

//when case#2 comes after case#1; problem solved

else p[x] ← double black

x ← p[x] //the same problem as at the beginning of case #2, just 1level above; case 2 **repeats**; in lgn problem solved

RB-delete- Case#3



- **Parent (B)=Red** (if after case #1) or **Black**, **brother (D)=Black**, **node C is Red**

- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees

(A=child of the deleted node, double black, pointed by x)




=> C<->D color interchange +right rotate => **case 4**

`brother[x] <- red`

`left[brother[x]] <- black`

`right_rotate(brother[x])`

RB-delete- Case#4



- **Parent (B)=Red** (if after case #1) or **Black**, **brother (D)= Black**, **node C is Red or Black**, **node E is Red**

- $\alpha, \beta, \gamma, \delta, \epsilon, \xi$ are RB trees

(\underline{A} =child of the deleted node, double_black, pointed by x)

=> $B \leftrightarrow D$ color interchange, $E \leftarrow \text{black} + \text{left rotate}$ => problem solved

`brother[x] ← color[parent[x]]`

`parent[x] ← black`

`right[brother[x]] ← black`

`left_rotate(p[x])` //1 more black node on x 's branch

RB-del – Rebalancing eval

- Case #1 rotation followed by any other case
 - $1+2 \Rightarrow$ problem solved $O(1)$
 - $1+3+4 \Rightarrow$ problem solved $O(1)$
 - $1+4 \Rightarrow$ problem solved $O(1)$
- Case #2 (**no rotation**, only recoloring) repeats 1 level up in the tree
 - Worst case $O(\lg n)$
 - Best case $O(1)$
- Case #3 rotation followed by case #4 $O(1)$
- Case #4 rotation; solves the problem $O(1)$
- Delete $O(\lg n)$ + rebalancing
 - Worst case: #2 repeats (recoloring only) $O(\lg n)$
 - Best case: #4 \Rightarrow 1 rotation $O(1)$
 - Other cases: #1+2 or 1+ 3+4 \Rightarrow 2 or 3 rotations $O(1)$
- $O(\lg n)$ overall, at most 3 rotations

RB-del - procedure

rb_del_fix(T,x)

while x<>root[T] and color[x]=black
do

if x=left[p[x]]

then

w<-right[p[x]]

if color[w]=red

then

color[w]<-black

color[p[x]]<-red

left_rotate(T,p[x])

w<-right[p[x]] //end case #1;



//another case comes

//cases on the left

//else case symmetric on the right; not discussed

//w=x's brother!

//case #1 APPLY ; coloring+rotation





RB-del - procedure

if color[left[w]] = black and color[right[w]] = black
then //case #2

color[w] ← red
x ← p[x]

else



RB-del – procedure - cont

else

//color[left[w]] ≠ black or color[right[w]] ≠ black

if color[right[w]] = black //E is black

then



//case #3

color[left[w]] ← black

color[w] ← red

right_rotate(T, w)

w ← right[p[x]] //end case #3



RB-del – procedure - cont

```
color[w] ← color[p[x]]
```

//case #4




```
color[p[x]] ← black
```

```
color[right[w]] ← black
```

```
left_rotate(T, p[x])
```

```
x ← root[T]
```

else ... //x=right[p[x], all 4 cases symmetric to the right
color[x] ← black



Conclusions on balanced search trees

Tree	Height	Ins	Del
BST	$[\lg n, n]$	$O(h)$	$O(h)$
RBT	$[\lg n, 2\lg n]$	2 rot	3 rot
AVL	$[\lg n, 1.45\lg n]$	1 rot	$\lg n$ rot
PBT	$\lg n$	n rot	n rot

For RBT, at most $\lg n/2$ color updates needed